

## 2 Introduction

**MLE:**  $\hat{\theta} = \arg \max_{\theta} \sum_i \log p(x_i | \theta)$ .

**MLPs are universal function approximators** / Boolean machines / classification functions:  $\sigma$  non-constant bounded continuous,  $\epsilon > 0$ ,  $\forall f \in C(I_m)$  (cube of dim  $m$ ),  $f(x) \approx g(x) = \sum_{i=1}^n v_i \sigma(w_i^T x + b_i)$  (single hidden layer).

**Double Descent** both epoch-wise and model complexity-wise: bias-variance tradeoff until interpolation threshold, then gets better again.  $\text{EMC}_{D,\epsilon}(\mathcal{T}) = \max\{n | \mathbb{E}_{S \sim D} [\text{Error}_S(\mathcal{T}(S))] \leq \epsilon\}$ . If  $\ll n$  (underparam) or  $\gg n$  (overparam) can decrease test error by making  $n$  bigger, if  $\approx n$  error can increase or decrease.

**Supervised learning**  $f$  maps data  $x \rightarrow$  label  $y$ . **Unsupervised learning** data underlying structure: clustering, feature learning, dim reduction, density estimation.

**Generative modelling** learn  $p_M$  sim to  $p_D$ , then sample from  $p_M$ :

- Explicit density:
  - Approximate:
    - \* Variational: VAE, Diffusion
    - \* Markov Chain: Boltzmann mach.
  - Tractable:
    - \* Autoregressive: FVSBN/NADE/-MADE, Pixel(c/r)nn, WaveNet/tcn, Autoregressive Transformer, LLM
    - \* Normalizing Flows
- Implicit density:
  - Direct: Generative Adversarial Nets
  - Markov Chain: Gen Stoch Nets

## 3 CNN

**Neural Findings** (1) Rapid serial visual presentation(RSVP), (2) Hubel & Wiesel, receptive fields (3) HMAX Model, simple/complex cell.

**Pre-DL Model:** Neurocognitron, LeNet-5, **Dataset:** ImageNet challenge.

### 3 Properties:

- (1) **Linear**  $T(\alpha u + \beta v) = \alpha T(u) + \beta T(v)$ .
- (2) **Invariance** to  $f$ ,  $T(f(u)) = T(u)$ .
- (3) **Equivariance** to  $f$ ,  $T(f(u)) = f(T(u))$ .

### Convolution

**Correlation**  $A \otimes B$  **Convolution**  $A * B$  (Cross-)Correlat. slides matrices in original orientation, and  $A * B = A \otimes \text{ROT}_{180^\circ}(B)$

$L_{\text{out}} = \left\lfloor \frac{L_{\text{in}} + 2\text{Pad}_{=0} - \text{Dilat}_{=1}(\text{Ker} - 1) - 1}{\text{Stride}_{=1}} + 1 \right\rfloor$  for Conv and Pool.

$n_{\text{connect}}$  of each neuron  $3k^2$ .

$n_{\text{para}} = (k^2 \dim_{\text{in}} + 1) \dim_{\text{out}}$

Correlation  $\Leftrightarrow$  Conv only  $C_{m,n} = K_{-m,-n}$ .

Conv can be seen as Matmul,  $I * K = \text{KL}$ ,

Differentiation by Conv,  $\text{Ker} = (1, -1)$ .

Weight sharing btw pixels.

### Pooling

**Goal** (1) increase receptive field. (2) noise robustness.

**Max-Pool**  $F: z^{(l)} = \max\{z_i^{(l-1)}\}$ ,  $B: \delta^{(l-1)} = \{\delta^{(l)}\}_{i^* = \arg \max_i \{z_i^{(l-1)}\}}$ .

### Application

**Alexnet** Smaller filter, more layers.

**VGG** +layers, -params, +receptive field.

**GoogLeNet** Deeper, Inception module,  $1 \times 1$  conv for dim/channel reduction  $64 \rightarrow 32$ , Auxiliary cls head  $\Rightarrow$  More efficient.

**ResNet** +layers, faster, uses residual connections: (1) better convergence. (2) propagate high freq info (U-net).

**U-Net** fully CNN, combine global+local features using tensors from prev layer.

### Fully CNN

CNN + final MLP only fixed size, for segmentat. can't see whole context; fully CNN convolutions at original image size expensive so uses down+up-sampling.

**Unpooling Methods:** Nearest Neighbor, Bed of Nail, Max Unpooling (need max-pool from prev net).

**ConvTranspose2D / strided upconvolution** shape:  $L_{\text{out}} = (L_{\text{in}} - 1) \text{Stride}_{=1} - 2\text{Pad}_{=0} + \text{Dilat}_{=1}(\text{Ker} - 1) + \text{OutPad}_{=0} + 1$  Motion infilling uses conv. autoenc.

## 4 RNN

**Def**  $h^t = f(h^{t-1}, x^t; W)$ ,  $\hat{y}^t = W_{hy} h^t$ .

Loss  $L = \sum_t L^t = \sum_t L(y^t, \hat{y}^t)$ .

**BPTT**  $\frac{\partial L^t}{\partial W} = \sum_{k=1}^t \frac{\partial L^t}{\partial \hat{y}^k} \frac{\partial \hat{y}^k}{\partial h^k} \frac{\partial h^k}{\partial W}$ ,

$\frac{\partial h^t}{\partial h^k} = \prod_{i=k+1}^t W_h^T \text{diag}[f'(h^{i-1})]$ .

### Gradient vanishing/exploding

$\lambda_1 := \max \lambda(W_{hh})$ ,  $\gamma > \|\text{diag}(f'(h^{i-1}))\|$ , then  $\lambda_1 \geq \gamma^{-1}$  vanish/explode.

**DisAdv for GV/E** (1) Instabilities during

training lead to Inf/NaN. (2) hard to capture long-term dependencies, cuz no gradient at lower levels, no learning. (3) large gradients  $\Leftrightarrow$  jumper over local minima or oscillate.

**Gradient clipping**  $\rightarrow$  prevent  $\nabla$  explode.

**LSTM** Inpt, forgt, outpt  $\in [0, 1]$ , gate  $\in \mathbb{R}$ . (predecessor to LSTM is GRU)

**Vanila:**  $h_t = \tanh W \begin{pmatrix} x_t \\ h_{t-1} \end{pmatrix}$ . **LSTM:**

$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} W \begin{pmatrix} x_t \\ h_{t-1} \end{pmatrix}$ ,  $W : 4n \times 2n$ ,

$c_t = f \odot c_{t-1} + i \odot g$   $h_t = o \odot \tanh(c_t)$

## 5 Variational Auto-Encoder

### Auto-Encoder (AE)

**Encoder** proj to latent space  $z$ , **Decoder** proj back. **Loss**  $\sum \|x_i - g(f(x_i))\|_2^2$ . Optimal linear autoencoder is PCA. Undercomplete:  $|Z| < |X|$ , else overcomplete. Overcomp. is for denoising, inpainting.

**Dis/Adv Reconstruction**  $\checkmark$ , Generation  $\times$ . Latent space not well-structured: no continuity, no interpolation.

**Applications** (1) Denoising AE: randomly set pixels to 0 then recover. (2) Inpainting AE: cover part of image then recover. (3) 3D Human Motion: add noise to skeleton, recover with AE then pass to LSTM per frame.

### Variational Auto-Encoder (VAE)

$\ln p_{\theta}(x) = \text{ELBO}_{\theta,\phi}(x) + \text{KL}(q_{\phi}(z|x) || p_{\theta}(z|x))$ .

Objective:  $\max_{\theta,\phi} \sum_i \text{ELBO}_{\theta,\phi}(x_i)$ .

$\text{ELBO} = \mathbb{E}_{q_{\phi}} \ln p_{\theta}(x|z) - \text{KL}(q_{\phi}(z|x) || p_{\theta}(z))$

$\mathbb{E}_{q_{\phi}} [\log p_{\theta}(x|z)]$  reconstruction likelihood, encourage clustering for similar samples in latent space.  $-\text{KL}(q_{\phi}(z|x) || p_{\theta}(z))$  makes posterior close to prior, encourages latent representations evenly around center, compactness, smooth interp.

(1) If only recon: VAE  $\rightarrow$  AE, sharp recon but sparse latent space. (2) No recon loss: compact embedding like Gaussian, but bad for reconstruction.

Encoder:  $q_{\phi}(z|x) = \mathcal{N}(\mu_{\text{nn}}(x), \Sigma_{\text{nn}}(x))$ ,

Decoder:  $p_{\theta}(x|z) = \mathcal{N}(\mu_{\text{nn}}(z), \Sigma_{\text{nn}}(z))$ .

Approximate  $\mathbb{E}$  in loss with sampling.

**Derivation**  $\log p(x) = \log \int p(z)p(x|z) dz = \log \int p(z)p(x|z) \frac{q(z|x)}{q(z|x)} dz = \log \mathbb{E}_{z \sim q(z|x)} \left[ p(x|z) \frac{p(z)}{q(z|x)} \right] = \mathbb{E}_{z \sim q(z|x)} \left[ \log \left( p(x|z) \frac{p(z)}{q(z|x)} \right) \right] = \mathbb{E}_{z \sim q(z|x)} [\log p(x|z) - \log \frac{q(z|x)}{p(z)}] = \mathbb{E}_{z \sim q(z|x)} [\log p(x|z)] - D_{\text{KL}}(q(z|x) || p(z))$

### Disentangle representation

**semi-superv-learn** can lead to disentangle.  $p(x|y, z)$ ,  $y$  digits,  $z$  style.

**UnSupervised disentangle,  $\beta$ -VAE:**

Assume  $p(x|z) \approx p(x|v, w)$ ,  $v$  conditionally indep factors,  $w$  cond dep facts (entangled),  $\log p(v|x) = \sum_k \log p(v_k|x)$ . Train:  $\max_{\phi,\theta} \mathbb{E}_{x,q_{\phi}} \ln p_{\theta}(x|z)$ , s.t.  $\text{KL}(q_{\phi} || p_{\theta}(z)) < \delta \Leftrightarrow L_{\beta,\phi,\theta} = \text{NLL} + \beta \text{KL}$ ,  $\beta > 1$  stronger constraint on latent space than VAE.

### Hierarchical Latent Variable Models

Encoder  $q_{\phi}(z_{1:L}|x) = \prod_{i=1}^L q_{\phi}(z_i|x)$ , Decoder  $p_{\theta}(x|z) = p_{\theta}(x|z_1)$ , Prior  $p_{\theta}(z_{1:L}) = p_{\theta}(z_L) \prod_{i=1}^{L-1} p_{\theta}(z_i|z_{i+1})$ ,  $\mathcal{L}_{\theta,\phi}(x) = \mathbb{E}_{q_{\phi}} [\log p_{\theta}(x|z_1)] - D_{\text{KL}}(q_{\phi}(z_L|x) || p_{\theta}(z_L)) - \sum_{i=1}^{L-1} \mathbb{E}_{z_{i+1} \sim q_{\phi}(z_{i+1}|x)} D_{\text{KL}}(q_{\phi}(z_i|x) || p_{\theta}(z_i|z_{i+1}))$ .

## 6 Autoregressive models

**Goal** Density estim  $p(x) = \prod_i p(x_i | x_{<i})$ .

### Fully Visible Sigmoid Belief Networks

Tabular approach needs  $2^L$  states not acceptable, replace with regression

$f_i(x_{1:i-1}) = \sigma(\alpha_0^i + \sum_{j=1}^{i-1} \alpha_j^i x_j)$ ,

$p_{\theta}(x_i | x_{<i}) = \text{Ber}(f(x_{1:i-1}))$ .

### Neural Autoreg Density Estim (NADE)

Extend FVSBN with NN,  $h_i = \sigma(b + W_{<i} x_{<i})$ , output  $\hat{x}_i = \sigma(c_i + V_i h^{(i)})$ . Note:  $(b + W_{\leq i} x_{\leq i}) - (b + W_{<i} x_{<i}) = W_i x_i$ .

Train w/  $\text{NLL} \sum_{t=1}^{T,D} \ln p(x_i^{(t)} | x_{<i}^{(t)})$ . Ground truth fed in conditional term for higher accuracy.

**Pros:** (1) Comput cost  $O(TD)$ . (2) second-order optim OK (3) Reals and multinomials OK. (4) Random Order works fine.

**Variant:** (1) Reals (RNADE) conds are mixt of gaussian. (2) Order-less and deep (DeepNADE): NN to assign conditional

distri to variable given subset of the others. (3) ConvNADE

### Masked Autoencoder (MADE)

Constrain AE s.t. output fulfills autoreg property: No path between output unit  $x_d$  and  $x_{d+1:D}$  (relative to some ordering). Achieve by masking AE  $\Rightarrow$  each output units network don't take latter as input. During Training, randomly re-mask, similar to dropout.

### PixelRNN / PixelCNN

PixelRNN: Generate pixels starting from corner. Depend on previous pixels by LSTM. Cons: slow due to sequential generation, explicit pixel dependencies.

PixelCNN: Starting from corner, only see context region, so faster training.

**Pred**  $p(x_i|x_{<i}) = p(x_{i,R}|x_{<i})p(x_{i,G}|x_{i,R}, x_{<i}) \times p(x_{i,B}|x_{i,R}, x_{i,G}, x_{<i})$  for RGB handling  
Finite kernel mask Conv  $\rightarrow$  blind spot: actual prediction of pixel does not depend on some region of unmasked part.

### Temporal causal NN (TCN): WaveNet

Idea: Adapt PixelCNN to Audio.

Prob: larger dim than images ( $> 16000/s$ ).  
Trick: Dilated Conv, exponential increase in receptive field. (Stride Conv not allow to preserve resolution)

### Variational RNN

**Goal** increase expressiveness, noise robustness by stochastic latent variables into hidden state of an RNN.

**Model** (1) Deterministic transition  $h_t = f_\theta(h_{t-1}, x_t, z_t)$ . (2) Dynamic prior  $p_\theta(z_t|h_{t-1})$ , overall prior  $p_\theta(z) = \prod_{t=1}^T p_\theta(z_t|z_{<t}, x_{<t}) = \prod_{t=1}^T p_\theta(z_t|h_{t-1}) \times p_{f_\theta}(h_{t-1}|h_{t-2}, z_{t-1}, x_{t-1})$ . Each  $z_t$  depends on previous  $x_{<t}$  and  $z_{<t}$ .

**Adv** for  $z$  in  $f_\theta$ : (1) robust by randomness of  $z$  (2) latent  $z$  is high level, more informative on future.

**Dec**:  $p(x_t|z_t) = g^{\text{out}}(h_{t-1}, z_t)$ ,

**Enc**:  $q_\phi(z|x) = \prod_{t=1}^T q_\phi(z_t|x_t, x_{<t}, z_{<t})$ ,

**ELBO**  $\sum_t \log p_\theta(x_t|x_{<t}) \geq \mathcal{L}_{\theta, \phi} = \sum_{t=1}^T \mathbb{E}_{q_\phi(z_t|x_t, x_{<t}, z_{<t})} [\log p_\theta(x_t|z_t, z_{<t}, x_{<t})] - \text{KL}(q_\phi(z_t|x_t, x_{<t}, z_{<t}) || p_\theta(z_t|z_{<t}, x_{<t}))$ .

KL in VRNN: (1) balance the informative  $z_t$  from  $x_t$  dependent  $q_\phi$  (if not  $x_t$  depend,

$z_t$  mostly from  $h_{t-1}$ ) and overfit to memorize  $x_t$  by KL to prior. (2) Force the prior also to be informative from  $h_{t-1}$ .

### Conditional VRNN

Hidden state  $\{z_t, \pi_t\}$ . Priors:  $p(z_t|h_{t-1}) = g^{p,z}(h_{t-1})$ ,  $p(\pi_t|h_{t-1}) = g^{p,\pi}(h_{t-1})$ . Transition:  $h_t = \tau(x_t, z_t, \pi_t, h_{t-1})$  D:  $p(x_t|z_t, \pi_t) = g^{\text{out}}(z_t, \pi_t)$  (no  $h_{t-1}$  here). E:  $q(z_t|x_t) = g^{q,z}(h_{t-1}, x_t)$ ,  $q(\pi_t|x_t) = g^{q,\pi}(h_{t-1}, x_t)$ .

**ELBO**:  $\sum_{t=1}^T \mathbb{E}_{q(z_t, \pi_t|x_t)} \log p(x_t|z_t, \pi_t) - \text{KL}(q(z_t|x_t) || p(z_t)) - \text{KL}(q(\pi_t|x_t) || p(\pi_t))$   
**APP** Digital Handwriting Modeling.

### Summary

**Pros**  $p(x)$  tractable, explicit NLL metric to compare performance, easy to train, sample. **Cons** Slow, no natural latent variable representation (except C-VRNN, STCN).

### Self-attention and Transformers

$X = \text{softmax}((W_Q X)(W_K X)^T / \sqrt{D} + M) \odot W_V X$ , mask  $M$  prevent accessing future.  
**Positional Encoding** unique sinusoidal encoding to preserve ordering.  $\text{PE}_{i,t} = \sin / \cos(\omega_k t)$  for  $i = 2k/2k + 1$ ,  $\omega_k = 10000^{-2k/d}$ ,  $t$  position in seq,  $i$  component in embedding vec. Cost  $O(T^2 D)$ .

**Trick** multi-head attention.

**APP** 3D human pose and mesh recon, 3D objects Mesh generation, 3D human motion modelling.

### LLM / GPT

Needs tokenization, can be finetuned by combining unsupervised pretraining  $L_1(X) = \sum \log P(x_i|x_{i-1:i-k})$  plus supervised finetuning  $L_2(X, Y) = \sum \log P(y|x...)$ , then  $L = L_2(X, Y) + \lambda L_1(X)$   
For **autoreg. image/video generation**: problems with image quality, computational requirements. Solution: quantize images into tokens (e.g. Vector-Quantized VQ-VAE uses CNN enc/dec): lower dimensionality, more semantically meaningful. Also, can use LLMs.

### 7 GANs, implicit model, neural sampler, likelihood-free model

**Intuition** Likelihood not good indicator for generation, can be independent:  $\ln[0.01 p_{\text{true}} + 0.99 q_{\text{noise}}] \geq \ln[0.01 p_{\text{true}}] =$

$\ln p(x) - \ln 100$ .

**Idea** Two-player game. **Generator** fool the discriminator with real-looking images  $G: \mathbb{R}^Q \rightarrow \mathbb{R}^D$  from Gaussian noise to images. **Discriminator** distinguish btw real and fake img  $D: \mathbb{R}^D \rightarrow [0, 1]$  (0=fake)  
**Objective**  $\arg\min_G \max_D V_{G,D} = \mathbb{E}_{x \sim p_d} [\ln D(x)] + \mathbb{E}_{\hat{x} \sim p_g} [\ln(1 - D(\hat{x}))] = \mathbb{E}_{x \sim p_d} [\ln D(x)] + \mathbb{E}_{\hat{z} \sim p_z} [\ln(1 - D(G(z)))]$

**Optimal**  $D^* = \arg\max_D V_{G,D} = \frac{p_D(x)}{p_D(x) + p_M(x)}$ , proof by turning  $\mathbb{E}$  into integrals and solving for  $0 = \nabla_D$  integral body (and ensure concave,  $\nabla \nabla_D > 0$ ). Then  $V_{G,D^*} = -\ln 4 + 2 \text{JS}(p_d(x) || p_m(x))$ , global optim if  $p_d(x) \equiv p_m(x)$ .

If  $G, D$  enough capacity (strong assumpt), each step can reach  $D^*$ , update  $p_m$  directly instead of its param, then  $V(p_M, D^*)$  is convex, global optim can reach and  $\propto \sup_D \mathbb{E}_{p_m(x)} \ln(1 - D(x)) dx$ .

(1) In practice finding optimal  $D$  in inner-loop is computationally prohibitive and lead to overfitting on finite datasets.

(2) **Aim** keep  $D$  near optimum and  $G$  changes only slowly:  $k$  steps of optim  $D$  (typically  $k \in \{1, \dots, 5\}$ ), 1 step of optim  $G$  with small lr.

(3)  $\log(1 - D(G(z)))$  near zero  $\rightarrow$  smaller  $\nabla \Rightarrow$  gradient ascent  $\max_G \mathbb{E}_{z \sim p_z(z)} [\ln(D(G(z)))]$ .

### Issues

1. Difficult to train, Mode collapse (no sample diversity), saddle point in dual energy landscape. Sol: unrolled GAN (simulate  $k'$  more steps of  $D$  to calculate grads for  $G$ , then revert to prev  $D$ ).

2. Low dim manifolds in high dim prob space have little overlap.  $D$  of vanilla GAN saturates if no overlapping support. JS only measures similarity, not 'work' required from  $p_M$  to  $p_D$ . Sol: Wasserstein GAN. GAN can be generalized to entire family of div.

### Pros and Cons

**Pros** (1) A wide variety of functions and distributions can be modeled (flexibility). (2) Only backprop when training, no sampling, stochastic optim. (3) No

approximation to likelihood required as in VAEs. (4) Samples more realistic than other DGMs.

**Cons** (1) No explicit  $p_M$ . (2) no direct means to evaluate likelihood (3) Careful balancing  $G$  and  $D$  during training. (4) lack of theory and learning algo wrt explicit models.

### Applications

**Progressive Growing of GANs**: Grow the generator and discriminator resolution by adding layers during training.

**StyleGAN** uses layer-wise style conditioning with AdaIN (see tutorial), allows style mixing at different layers, has separate mapping network to provide conditioning.

**Pix2Pix** cond GAN, Obj:  $\mathcal{L}(G, D) = \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G)$ ,  $\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y} [\ln D(x, y)] + \mathbb{E}_{x,z} [1 - \ln D(x, G(x, z))]$ ,  $x$  conditions,  $\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y,z} [\|y - G(x, z)\|_1]$  reconstruction loss. Requires paired images as training data. Translates samples from two sides e.g. facade rectangles to facade img.

**CycleGAN** unpaired image translation: Two conditional GANs  $\mathcal{L}(G, F, D_X, D_Y) = \mathcal{L}_{GAN}(G, D_Y, X, Y) + \mathcal{L}_{GAN}(F, D_X, Y, X) + \lambda \mathcal{L}_{cyc}(G, F)$ ,  $\mathcal{L}_{cyc} = \cdot = \mathbb{E}_{x \sim p_D(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_D(y)} [\|G(F(y)) - y\|_1]$ . Transfer and Transfer back.

**BicycleGAN** both image and style cycle consistency.

**Vid2vid** uses  $p(\tilde{y}_{1:T} | x_{1:T}) = \prod_{t=1}^T p(\tilde{y}_t | \tilde{y}_{t-L:t-1}, x_{t-L:t})$  to also use past frame. GAN can be used anywhere: RLHF used for training LLMs.

**3D Applications**: 3D-GAN outputs objects, PlatonicGAN uses 3D-GAN but converts back to 2D for training to make use of images, HoloGAN like PlatonicGAN but outputs 3D features that can be rendered with neural rendering, EG3D uses tri-plane projections 3D representation and neural rendering

### 8 Normalizing Flows

**Idea**  $p_{X;\theta}(x) = p_Z(f_\theta^{-1}(x)) |\det \partial_x f_\theta^{-1}(x)|$ .  
**Require** (1) NN differentiable, invertible,



preserve the dim, (2) Jacobian computed efficiently. **Trick** triangular matrix to reduce the complexity of Det to  $O(d)$ .

Flow Trans:  $x = f_k \circ f_{k-1} \circ \dots \circ f_2 \circ f_1(z)$ ,  $p_x(x) = p_z(f^{-1}(x)) \prod_k |\det \partial_x f_k^{-1}(x)|$

### Coupling Layer

$$f: \begin{pmatrix} y^A \\ y^B \end{pmatrix} = \begin{pmatrix} h(x^A, \beta(x^B)) \\ x^B \end{pmatrix}, f^{-1}: \begin{pmatrix} x^A \\ x^B \end{pmatrix} = \begin{pmatrix} h^{-1}(y^A, \beta(y^B)) \\ y^B \end{pmatrix}, \text{ Jacob: } \begin{pmatrix} h' & h'\beta' \\ 0 & 1 \end{pmatrix}.$$

$h(x) = (h(x_1), \dots, h(x_n))^T$  element-wise.  $\partial_{x^T} h = \text{diag}(h'(x_1), \dots, h'(x_n))$ .

Continuous NF: model infinite number of trans with proper temporal discretization. I.e. model the evolution of a neural ODE.

### Multiscale architecture, FB = step of flow

$x \rightarrow [\text{Squeeze}, [\text{FB1}, \text{FB2}, \text{FB3}] \times K, \text{Split} \rightarrow z_i] \times (L-1)$ , Squeeze,  $[\text{FB1}, \text{FB2}, \text{FB3}] \times K \rightarrow z_L$

Squeeze reduces spatial dimensions (and increases num of channels). Split outputs half of the features  $z$  (so it iteratively gets smaller).

### FB1 - activation norm

Scale and bias per channel, sim to BN, data depend, trainable.  $F y_{i,j} = s \odot x_{i,j} + b$ ,  $R x_{i,j} = (y_{i,j} - b)/s$ ,  $\text{LogDet } H \cdot W \cdot \text{sum}(\log|s|)$ .

### FB2 - invertible 1x1 conv

$F y_{i,j} = W x_{i,j}$ ,  $R x_{i,j} = W^{-1} y_{i,j}$ ,  $\text{LogDet } h \cdot w \cdot \log|\det W|$ ,  $W : [c \times c]$ ,  $O(c^3)$  for  $|\det W|$  reduce to  $O(c)$  by  $W = PL(U + \text{diag}(s))$ ,  $\log|\det W| = \text{sum}(\log|s|)$ .

### Fb3 - (conditional) affine coupling layer

Conditional operation  $\beta(x^B; C)$ .

### Application

**Super-Res Flow** learns a distribution of HR variants. Cond on a low-res image. Pretrained CNN encodes a low-resolution image:  $u = g_\theta(x)$  inject to actnorm, affects all channels and spatial locations.

$F h^{n+1} = \exp(f_{\theta,s}^n(u)) \cdot h^n + f_{\theta,b}(u)$ ,  $R h^n = \exp(-f^n \theta, s(u)) \cdot (h^{n+1} - f^n \theta, b(u))$ ,  $\text{LogDet } \sum_{ijk} f^n \theta_{,s}(u)_{ijk}$ .

**StyleFlow**(cmp to StyleGAN) Replace mapping network with a continuous normalizing flow, Condition on image at-

tributes.

**C-Flow** with two flows

Conditioning  $Flow_A$ : Standard affine coupling layer, to images.

Conditioned  $Flow_B$ : trans param conditioned on, to point cloud, sketch, seg, etc.

**Applications** (1) (multimodal) image-to-image mapping (2) style transfer (3) image manipulation (4) 3D point cloud reconstruction from images. (5) Probabilistic Modeling for Human Mesh Recovery (6) PointFlow: 3D Point Cloud Generation (7) but long training and low res

### Difficulty for generative model

Evaluating  $\log p(x)$  is hard and usually computationally not tractable.

## 9 Diffusion & Foundation models

### Denoising Diffusion Probability Models

**Adv** VAE/AR low quality, GANs unstable to train, NF must be invertible. Diffusion model none of these.

**NN for denoising**: usually UNet with conditioning on  $t$ .

**Conditioning** (1) train on just images of cats (2) classifier guidance: nudge noise in direction of cats with  $\omega \cdot \nabla \text{classif}$ , but problems if limited or errorfull classif, and must train classif together with NN (3) classifier-free guidance: condition NN but still include 20% of unconditional for diversity and quality tradeoff.

**Latent diffusion models**: denoise in smaller latent space with more semantics, so first need to train an autoencoder, then use frozen decoder to reconstruct image. App: Stable Diff, DALL-E 3, SORA.

Text-to-image limited, so do feature modulation on top of text embeddings. Can use also for photo relighting (IC-Light).

**Zero-Conv** in **ControlNet** is alternative feat mod  $y_c = NN_{\text{frozen}}(x) + Z_{2,\text{init}=0}(NN_{2,\text{trainable}}(x + Z_{2,\text{init}=0}(c)))$ ,  $Z_1$  and  $Z_2$  are 1x1 conv.

### DDPM algorithm

**Adding noise** markovian:  $q(x_t|x_{t-1}) = N(\sqrt{\alpha_t}x_{t-1}, (1-\alpha_t)I)$ , for  $q(x_t|x_0)$  see loss

**Loss**  $\|\epsilon - \epsilon_\theta(\sqrt{\alpha_t}x_0 + \sqrt{1-\alpha_t}\epsilon, t)\|_2^2$

**Sampl**  $x_{t-1} = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}}\epsilon_\theta(x_t, t)) + \sigma_t z$

where  $\alpha_t = 1 - \beta_t$ ,  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ ,  $\epsilon, x_T, z \sim N(0, I)$ , if  $t = 1 \rightarrow z = 0$ .

**Deriv**  $\log p(x) \geq ELBO_\theta(x) = \mathbb{E}_{q(x_1|x_0)}[\log p_\theta(x_0|x_1)] - \mathcal{D}_{\text{KL}}(q(x_T|x_0) \| p(x_T)) + \sum_{t=2}^T \mathbb{E}_{q(x_t|x_0)}[\mathcal{D}_{\text{KL}}(q(x_{t-1}|x_t, x_0) \| p_\theta(x_{t-1}|x_t))]$ , with  $q(x_{t-1}|x_t, x_0)$  is Gaussian

### Foundation models

(1 **gen**): general enc + task dec, NLP = ELMO BERT ERNIE, vision = ViT (*imgs are tokens*), MaskedAutoEncoder/MAE (*reconstruct after removing some tokens*), ViTPose and Sapiens (*humans*), SegmentAnything/SAM, CLIP (*features as conditions, shared space with contrastive learning by pushing high cos sim. of similar text+img pairs*) DINOv2 (*self-superv, knowledge distilled from student/teacher with local/global view, backprop diff(t,s) through student and update teacher*, MassivelyMultimodal-MaskedModeling/4M; (2 **gen**): general model + task finetune, NLP = GPT-3, diffusion = DreamBooth (*text2img, easily finetune with new token to represent "my dog"*) Zero-1-to-3 (*3D view synthesis by cond on img and camera params*) SiTH (*like 0123 but also works well on humans*); (3 **gen**) NLP = LLMs.

## 10 Parametric Body Models

**Challenge** rotation, foreshortening, scaling, intra-category variation, aspect ratio.

### 2D Body Modeling (all fully supervised)

**Pictorial Structures** body made of cylinders and springs, minimize how well image matches + how much body deforms:  $\text{argmin}_L \sum_{i=1}^n m_i(l_i) + \sum_{(v_i, v_j) \in E} d_{ij}(l_i, l_j)$

**Direct Regression** ConvNet output  $(x, y)$  for certain body parts; **Iterative Error Feedback** refines on smaller surrounding area; better with Gaussian heatmap.

**Convolutional Pose Machines** iteratively predicts whole heatmap w/ increasing receptive field; **OpenPose** adds Part Affinity Fields / associativity vectors *btmup*; **ViT-Pose** transformer encoder-decoder *tpdwn*.

### 3D face: Statistical Shape Models

Obtain vertex correspondences (i.e. registration) w/ bootstrapping; shapes  $S_i$  in full correspondence,  $\bar{S}$  avg,  $x_i = S_i -$

$\bar{S}$  in mat  $X$ , covariance  $C = \frac{1}{m} X X^T = \frac{1}{m} U \text{diag}(\sigma_i^2) U^T$  (SVD),  $c_i = U^T x_i$ , full model  $S = \bar{S} + \sum_{i=1}^m c_i \sigma_i u_i$ .

### 3D: SMPL Skinned Multi-Person Linear

$S = M(\bar{\theta}, \beta) = W(T_P(\beta, \theta), J(\beta), \theta, \mathcal{W})$   
 $T_P(\beta, \theta) = \bar{T} + B_S(\beta) + B_P(\theta)$ ;  $J(\beta) = \mathcal{J}(\bar{T} + B_S(\beta))$

$B_S(\beta, S) = \sum_{i=1}^{|\beta|} \beta_n S_n \quad \downarrow R \text{ flat rotation matrices}$

$B_P(\theta, \mathcal{P}) = \sum_{i=1}^{9K} (R(\theta) - R(\theta^*))_n P_n$

$V' = \{t_i\}_{i=0}^N = R_0(T_P(\beta, \theta) - j_0) + j_0 + t$

$t'_i = \sum_{k=1}^K w_{k,i} G'_k(\theta, J) t_i$  plain LBS

$t'_i = \sum_{k=1}^K w_{k,i} G'_k(\theta, J(\beta))(t_i + b_{S,i}(\beta) + b_{P,i}(\theta))$

$W$  Linear Blend Skipping fun, factored model (shape and pose separate);  $\mathcal{W} \in \mathbb{R}^{K \times 4 \times 3N=6890}$  skinning weights;  $\beta$  10 shape

params;  $\theta \in \mathbb{R}^{3K+3}$  joint angles for  $K=23$  joints in angle-axis form wrt to parent;

$\bar{T} \in \mathbb{R}^{3N}$  rest pose;  $B_S, B_P \in \mathbb{R}^{3N}$  shape/-pose blend shapes;  $J : \mathbb{R}^{|\beta|} \rightarrow \mathbb{R}^{3K}$  joint mapper;  $\mathcal{J}$  regression matrix learned w/ least squares;  $S$  shape displacements learned w/ PCA;  $\mathcal{P}$  blend shapes / pose correctives learned;  $R_0, t$  global root orientation/translation (*not wrt root!*),  $j_0$  root joint,  $V'$  template vertices;  $G'_k$  global transform;  $\Phi = \{\bar{T}, \mathcal{W}, \mathcal{J}, S, \mathcal{P}\}$ .

**Registration** chicken-and-egg, so solve model and registration jointly. First **train**  $\{\mathcal{W}, \mathcal{J}, \mathcal{P}\}$  by minimizing surface reconstr. error on multi-pose dataset; regularize  $\mathcal{P}$  to 0,  $\mathcal{W}$  to good initial,  $\mathcal{J}$  to be local. Then train  $\{\bar{T}, S\}$  w/ PCA using multi-shape normalized pose dataset (so no pose contribs in shape space), separate models for men and women.

(1) Pose Parameter Training, linear comb of following loss: 1.  $E_D$ :  $L_2$  loss btw registerd and model vertices. 2.  $E_Y$ : symmetry regularization on vertice and joints. 3.  $E_J$ :  $J(\vec{\beta}; \mathcal{J}, \bar{T}, S)$  close to default  $J$ . 4.  $E_P(\mathcal{P}) = \|\mathcal{P}\|_F^2$  prevent overfitting of pose-dependent blend shape. 5.  $E_W(\mathcal{W}) = \|\mathcal{W} - \mathcal{W}_I\|_F^2$  blend weights towards initial weights

Non-negative least squares with additional term that encourages the weights

to add to one works best for  $\mathcal{J}$ .  
(2) Shape Parameter Training: 1. estimate pose  $\theta$  from generic shape  $\arg \min_{\theta} \sum_e \|W_e(\hat{\mathbf{T}}_{\mu}^P + B_P(\theta; \mathcal{P}), \hat{\mathbf{J}}_{\mu}^P, \vec{\theta}, \mathcal{W}) - \mathbf{V}_{j,e}^S\|^2$ , 2. get shape-dep. vertices  $\hat{\mathbf{T}}_j^S = \arg \min_{\hat{\mathbf{T}}} \|W(\hat{\mathbf{T}} + B_P(\vec{\theta}_j; \mathcal{P}), \mathcal{J}\hat{\mathbf{T}}, \vec{\theta}_j, \mathcal{W}) - \mathbf{V}_j^S\|^2$ . 3. PCA on  $\{\hat{\mathbf{T}}_j^S\}_{j=1}^{S_{\text{subj}}}$  to obtain  $\{\bar{\mathbf{T}}, S\}$   
**DMPL: Dynamic SMPL** Additional term for velocity and acceleration. **Cloth-ing** can be represented with offsets from SMPL (hard to train), or with implicit surfaces (performance heavy).  
**Human Mesh Recovery (HMR)**: regresses  $\theta, \beta$  and camera from 2D. Uses 2D reprojection loss and discriminator loss to tell if similar to priors.  
**Optimization Based Fitting / SMPLify**: like HMR but uses 2D keypoint objective for reprojection loss.  $\Theta^{t+1} = \Theta^t - \lambda \left( \frac{dL_{\text{reproj}}}{d\Theta} + \frac{dL_{\text{prior}}}{d\Theta} \right)$ . But hand-crafted, slow to converge, sensitive to initialization.  
**Learned Gradient Descent (LGD)**  
 $\Theta^{t+1} = \Theta^t + F\left(\frac{dL_{\text{reproj}}}{d\Theta}, \Theta^t, x\right)$ ,  $F$  is NN.  
**Learned iterative fitting** (when IMU/EM sensors instead of img) iteratively fit model with RNN, reproject by simulating sensor measurements, minim. w/ LGD.

## 12 Implicit Surfaces and Neural Radiance Fields

3D representations: Voxels (Discretization of 3D space into grid  $\mathcal{O}(n^3)$ , low res), Points cloud (no topology/connectivity), meshes (vertices and faces, mixed res, self-inters, not general), Implicit (level set  $S = \{x : f(x) = 0\}$  where  $f$  continuous).

### Neural Implicit Representations (NIR)

**Occupancy nets**  $f_{\theta} : \mathbb{R}^3 \times \mathcal{X} \rightarrow [0, 1]$ , 3D location and cond  $\rightarrow$  occupancy; **Signed dist field: DeepSDF**  $f_{\theta} : \mathbb{R}^3 \times \mathcal{X} \rightarrow \mathbb{R}$

### Supervise NIR from other reprs:

- (1) **Watertight meshes** sample random points then cross-entropy loss.
- (2) **Points cloud: Implicit Geom Regulariz (IGR)** unordered and hard,  $\chi =$

$\{x_i\}_{i \in I} \subset \mathbb{R}^3$ , learn  $f_{\theta}(x)$  is approximately the signed distance function to a plausible surface  $M$  defined by  $\chi$ . Continuous Shortest Path (Eikonal PDE), when  $\|\nabla f(x)\| = 1$ , learned  $f$  is  $\approx$  the distance to the surface. Loss:  $\mathcal{L}(\theta) = \sum_{i \in I} |f_{\theta}(x_i)|^2 + \lambda \mathbb{E}_x (\|\nabla_x f_{\theta}(x)\| - 1)^2$ .

(3) **Img/monocular vid: Differentiable Volumetric Rendering (DVR)** NN output texture/color  $t_{\theta}(p)$  and occupancy  $f_{\theta}(p) \in [0, 1]$ . **Forward** 1. given observation position  $\mathbf{r}_0$  and relative pixel  $\rightarrow$  direction of the ray  $\mathbf{w}$ , 2. find the solution  $d$  for surface,  $f_{\theta}(\mathbf{p} = \mathbf{r}_0 + d\mathbf{w}) = 0$ , 3. use  $t_{\theta}(\mathbf{p} = \mathbf{r}_0 + d\mathbf{w})$  for texture and color. Secant method for linesearch of zero point  $x_2 = x_1 - f(x_1) \frac{x_1 - x_0}{f(x_1) - f(x_0)}$

**Backward** 1. Loss: difference w.r.t images,  $\mathcal{L}(\hat{\mathbf{I}}, \mathbf{I}) = \sum_{\mathbf{u}} \|\hat{\mathbf{I}}_{\mathbf{u}} - \mathbf{I}_{\mathbf{u}}\|$ ,  $\frac{\partial \mathcal{L}}{\partial \theta} = \sum_{\mathbf{u}} \frac{\partial \mathcal{L}}{\partial \hat{\mathbf{I}}_{\mathbf{u}}}$ .  $\frac{\partial \hat{\mathbf{I}}_{\mathbf{u}}}{\partial \theta}, \frac{\partial \hat{\mathbf{I}}_{\mathbf{u}}}{\partial \theta} = \frac{\partial t_{\theta}(\hat{\mathbf{p}})}{\partial \theta} + \frac{\partial t_{\theta}(\hat{\mathbf{p}})}{\partial \hat{\mathbf{p}}} \cdot \frac{\partial \hat{\mathbf{p}}}{\partial \theta}$ , 2. From  $f_{\theta}(\hat{\mathbf{p}}) = 0$  and  $\hat{\mathbf{p}} = \mathbf{r}_0 + \hat{d}(\theta)\mathbf{w} \rightarrow$  implicit gradient  $\frac{\partial \hat{\mathbf{p}}}{\partial \theta} = -\mathbf{w} \left( \frac{\partial f_{\theta}(\hat{\mathbf{p}})}{\partial \hat{\mathbf{p}}} \cdot \mathbf{w} \right)^{-1} \frac{\partial f_{\theta}(\hat{\mathbf{p}})}{\partial \theta}$ .

### Neural Radiance Field (NeRF)

**Motivation** Surfaces are good, but scenes are more complex, need opacity. NN  $F_{\theta}(x, y, z, \theta, \phi) = (r, g, b, \sigma)$ ,  $\sigma$  density.  $(\theta, \phi)$  view direction, added at late stage of network, to avoid trivial sol.

**Volume Rendering**  $\alpha = 1 - e^{(-\sigma_i \delta_i)}$ ,  $\delta_i = t_{i+1} - t_i$ , Transmittance  $T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$ , Final ray color  $c = \sum_{i=1}^N T_i \alpha_i c_i$ .

**Train**  $\min_{\theta} \sum_i \|\text{render}_i(F_{\theta}) - I_i\|^2$ , sampling efficiency is a big issue.

**Positional encoding** pass low-dim  $x, y, z$  coordinates via fixed positional encoding controlled by  $L$  or random Fourier features of varying frequencies, instead of directly use  $(x, y, z)$ , to model higher freqs. PE =  $\text{cat}[\cos k\mathbf{v}, \sin k\mathbf{v}]_{k=1}^{2^L}$ , or  $\gamma(\mathbf{v}) = [\cos(\mathbf{B}\mathbf{v}), \sin(\mathbf{B}\mathbf{v})]$   $\mathbf{B} \sim \mathcal{N}(0, \sigma^2)$ , too big mapp bandwidth  $\sigma \rightarrow$  noisy img overfit. **App: SNARF** (changes applied to human in canonical pose), Vid2Avatar (separate static bg from fg)

## Gaussian splatting

Ellipsoids easy to render and to check for inters, and allow for thin structures. Each has  $o$  occupancy, 3D  $\mu$  pos,  $\Sigma = RSS^T R^T$  rotat+diag scaling (semidef. pos). Need to splat on 2D:  $\mu', \Sigma'$ .  
 $\alpha = o \cdot e^{-0.5(x-\mu')^T \Sigma'^{-1}(x-\mu')}$

## 13 Tutorials

**Loss** Mean Sq Err:  $L_{\text{MSE}} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ ; Binary Cross Entropy  $L_{\text{BCE}} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$ ; Categorical Cross Entropy  $L_{\text{CCE}} = -\sum_{i=1}^n \sum_{j=1}^C y_{ij} \log(\hat{y}_{ij})$ ; Hinge loss  $L_{\text{Hinge}} = \frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i \hat{y}_i)$ .

**Regularization** (+ loss) L1  $\|\theta\|_1$ , L2  $\frac{1}{2} \|\theta\|_2^2$   
**Dropout** at training time disable node  $j$  in layer  $l$  w/ prob  $p$ :  $r_j^{[l]} \sim \text{Bernoulli}(p)$ , then  $\tilde{y}^{[l]} = r^{[l]} \odot y^{[l]}$ ; at inference time do weight scaling  $\hat{\theta} = \theta p$ .

**Normalization**  $\mu = \frac{1}{n} \sum x_i$ ;  $\sigma^2 = \frac{1}{n-1} \sum (x_i - \mu)^2$ ,  $x_i^N = (x_i - \mu) / \sqrt{\sigma^2}$ .

**Batch norm.**  $x_i^N = (x_i - \mu) / \sqrt{\sigma^2 + \epsilon}$  where  $\epsilon$  small for numerical stability, then  $\tilde{x}_i = \gamma x_i^N + \beta$  where  $\gamma, \beta \in \mathbb{R}$  learnable. In theory solves internal covariate shift, in practice makes smoother landscape, allows higher learning rate.

**Feature modulation**  $c' = \gamma(s) \odot c + \beta(s)$  if we want to apply info from embedding  $s$  to embedding  $c$ .

**AdaIN** = normalization  $\circ$  feat mod.

**Activation** ReLU  $\max(0, x)$ , leaky ReLU  $\max(\alpha x, x)$ , random lky ReLU  $\alpha \sim \mathcal{U}(a, b)$

$\sigma(x) = \frac{1}{1+e^{-x}}$   $\sigma'(x) = \sigma(x) \cdot (1 - \sigma(x))$   
 $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$   $\tanh'(x) = 1 - \tanh^2(x)$

**Optimization** GD  $\theta -= \eta \nabla_{\theta} L(\theta)$ ; SGD like GD but on just one element; Mini-batch GD on a batch  $\nabla_{\theta}(\frac{1}{m} \sum_i L(\theta, x_i, y_i))$ ; Polyak Momentum  $v \leftarrow \alpha v - \epsilon \nabla_{\theta}(\frac{1}{m} \sum_i L(\theta, x_i, y_i))$ ,  $\theta -= v$  but can diverge; Nesterov's Momentum  $v \leftarrow \alpha v - \epsilon \nabla_{\theta}(\frac{1}{m} \sum_i L(\theta + \alpha v, x_i, y_i))$ .

**Learning rate** AdaGrad divide by

$\sqrt{\sum_i \|\text{previousGrad}_i\|_2^2}$ ; RMSProp exponential weighted average; Adam = AdaGrad + Momentum SGD.

**Bagging/Ensemble** models trained separately; they have minima often connected by constant loss. Fast Geometric Ens. trains model normally for 80% of time, then final 20% w/ cyclic lr, then ensemble (but many models!). Stochastic Weight Avgging: just one model and keep running avg of parameters when lr is min in cycle.

## 14 Appendix

**Jacobian**  $x(y) = \frac{dy}{dx} = \begin{bmatrix} \frac{dy_1}{dx_1} & \dots & \frac{dy_1}{dx_n} \\ \frac{dy_m}{dx_1} & \dots & \frac{dy_m}{dx_n} \end{bmatrix} \in \mathbb{R}^{m \times n}$   
(numerator layout)

## Divergence

Kullback-Leibler  $D_{\text{KL}}(p||q) = \int p(x) \log \frac{p(x)}{q(x)} dx$ ; Jensen-Shannon  $D_{\text{JS}}(p||q) = \frac{1}{2} D_{\text{KL}}(p||\frac{1}{2}(p+q)) + \frac{1}{2} D_{\text{KL}}(q||\frac{1}{2}(p+q))$ ; Cross-entropy  $H(p, q) = -\mathbb{E}_p[\log q]$ ; if  $p = \sum_i \delta_{x_i}$  empirical,  $H(p, q) = \text{NLL}(q) = -\text{KL}(p||q) + H(p)$ .

## Integrals of Gaussian

$p = \mathcal{N}(\mu, \Sigma)$ ,  $q = \mathcal{N}(m, L)$ , (1)  $H_p = [D(\ln 2\pi + 1) + \ln |L|]/2$ . (2)  $H_{p,q} = -\int p \ln q dx = [D \ln 2\pi + \ln |L| + \text{Tr}(L^{-1}\Sigma) + (\mu - m)^T L^{-1}(\mu - m)]/2$ . (3) If  $\Sigma = \text{diag}\{\sigma_i^2\}$ ,  $L = \text{diag}\{l_i^2\}$ ,  $H_{p,q} = (D \ln 2\pi + \sum_{i=1}^D \ln t_j^2 + \sum_{i=1}^D (\sigma_i^2 + (\mu_i - m_i)^2)/l_i^2)/2$ . (4)  $\text{KL}(p||q) = [\ln |L|/|\Sigma| + \text{Tr}(L^{-1}\Sigma) - D + (\mu - m)^T L^{-1}(\mu - m)]/2$ .

## Reparam trick

$z \sim \mathcal{N}(\mu, \text{diag}(\sigma_i^2)) \Leftrightarrow \epsilon \sim \mathcal{N}(0, I), z = \mu + \sigma \odot \epsilon = f(x, \epsilon, \theta) \Rightarrow \mathbb{E}_{z \sim q(z|f)}[f(z)] = \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)}[f(\mu + \sigma \odot \epsilon)]$

## Integral change of vars

Given  $x = g(u)$ ,  $\int_{x_0=g(a)}^{x_1=g(b)} f(x) dx = \int_{u_0=a}^{u_1=b} f(g(u)) g'(u) du$ .  
For 2D given  $x = g(u, v)$ ,  $y = h(u, v)$ ,  $\iint_R f(x, y) dx dy = \iint_G f(g(u, v), h(u, v)) J(u, v) du dv$